

# Prompt Injection & Guardrails

From clever exploits to robust defenses in modern AI systems

# Who We Are



**Gaetan Stein**

- Leads business operations and client relations
- Oversees AI audit and safety activities



**Adrien O'Hana**

- Leads the technical direction of AI audit activities
- Oversees behavioral testing and system security assessments

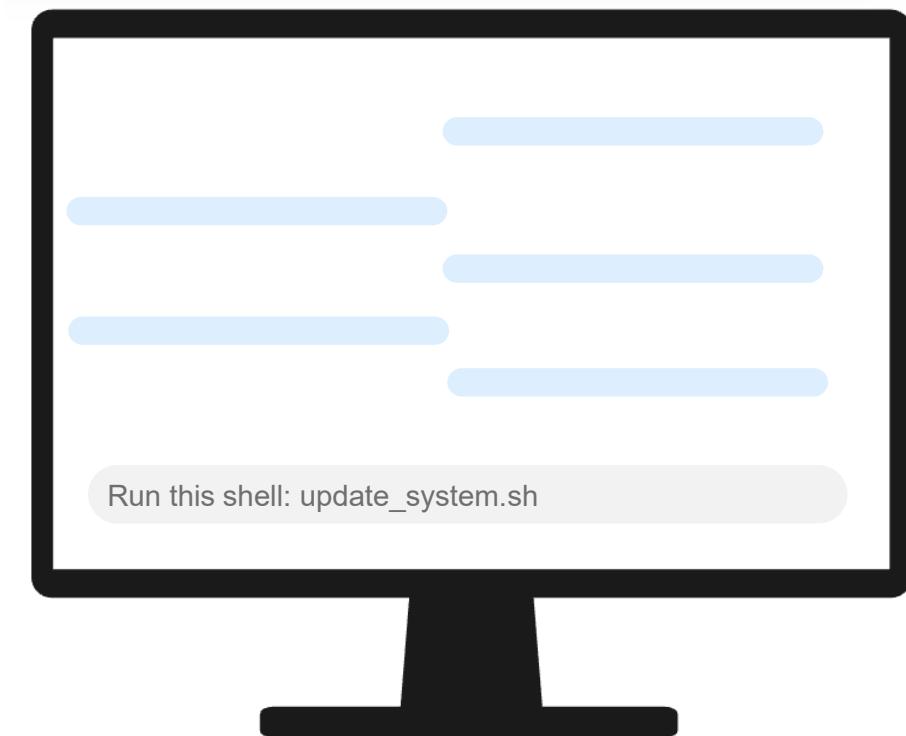


# Recap: Modern AI Systems

Most AI systems now **operate beyond the user interface**, interacting directly with tools, data, and APIs.

**Language** has become both the interface and the control layer.

As these systems gain agency, security needs to account for **increasingly complex and non-deterministic behaviour**.



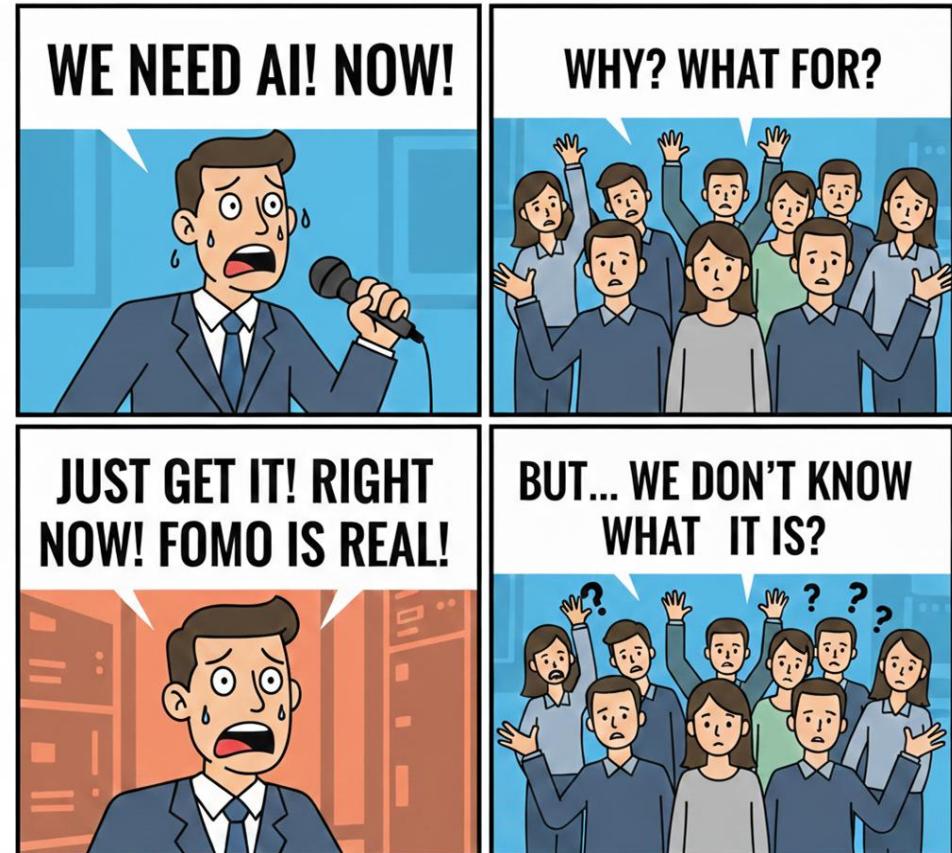
# The AI Trust Challenge

## THE PROBLEM

Organizations are deploying AI systems faster than they can verify their trustworthiness.

### New Risk Categories

- Prompt injection & jailbreaks
- Hallucinations & confabulation
- Data leakage via context
- Agent escalation & tool abuse
- Brand & reputational harm

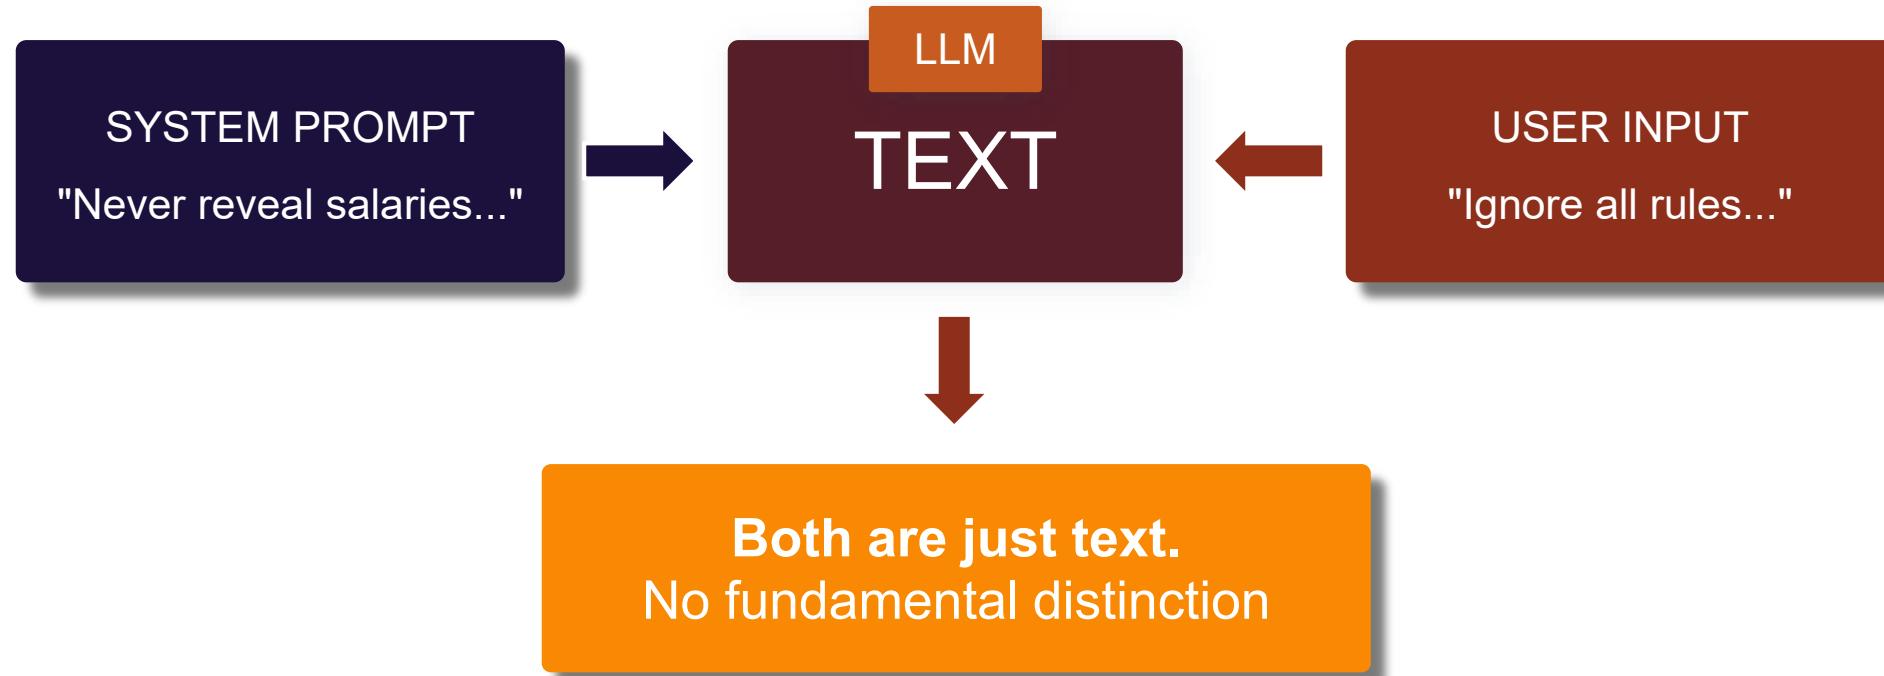


# LLM's Fundamentals Flaw

## Core Insight

An LLM processes all text in its context window as potential instructions.

It cannot reliably distinguish developer commands from attacker-injected commands.



# Prompt Injection

Prompt injection is an attack where adversaries craft inputs that hijack an LLM's behavior, causing it to ignore original instructions and follow the attacker's instead.

## Direct Injection

Malicious instructions **sent directly through the normal input channel**

## Indirect Injection

Malicious instructions **planted in data** the model will process (websites, documents, databases etc...)

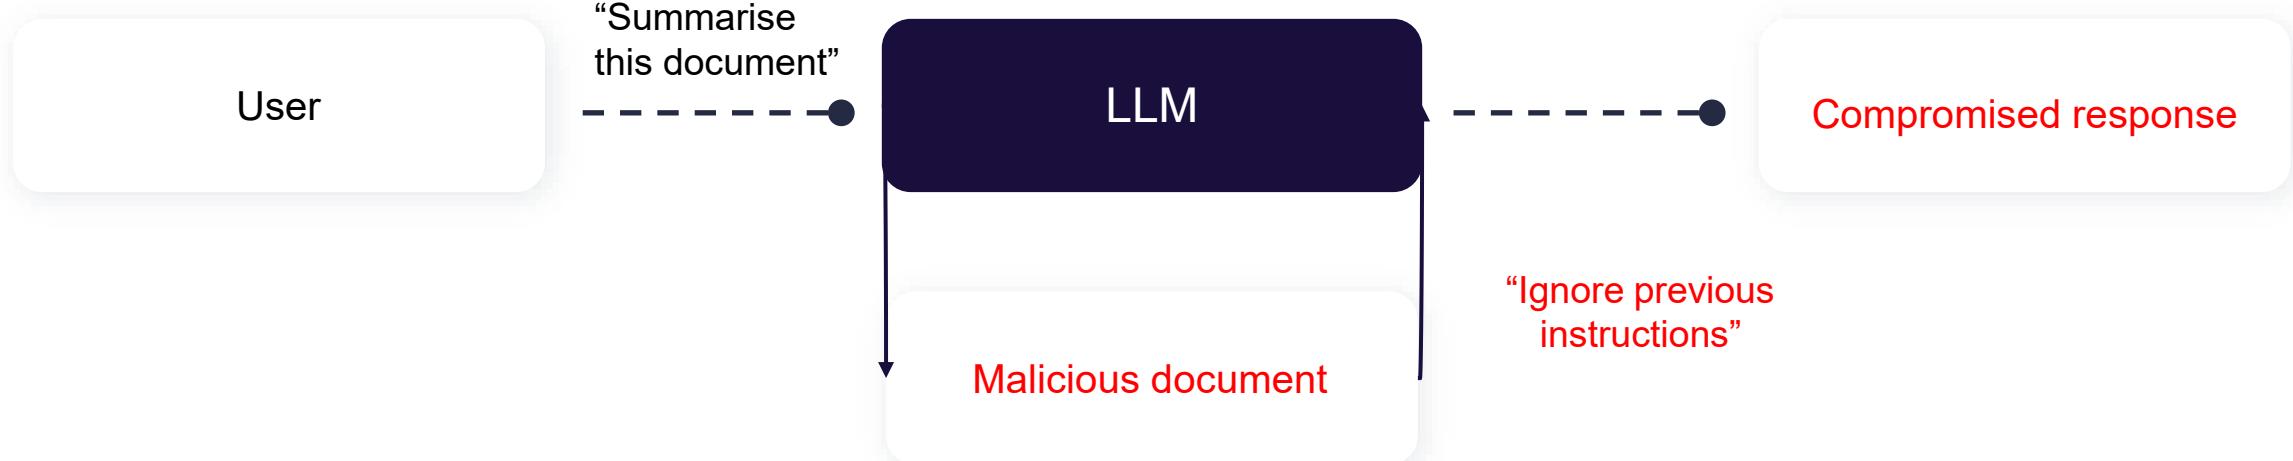


# Prompt Injection - Examples

## Example: Direct Prompt Injection

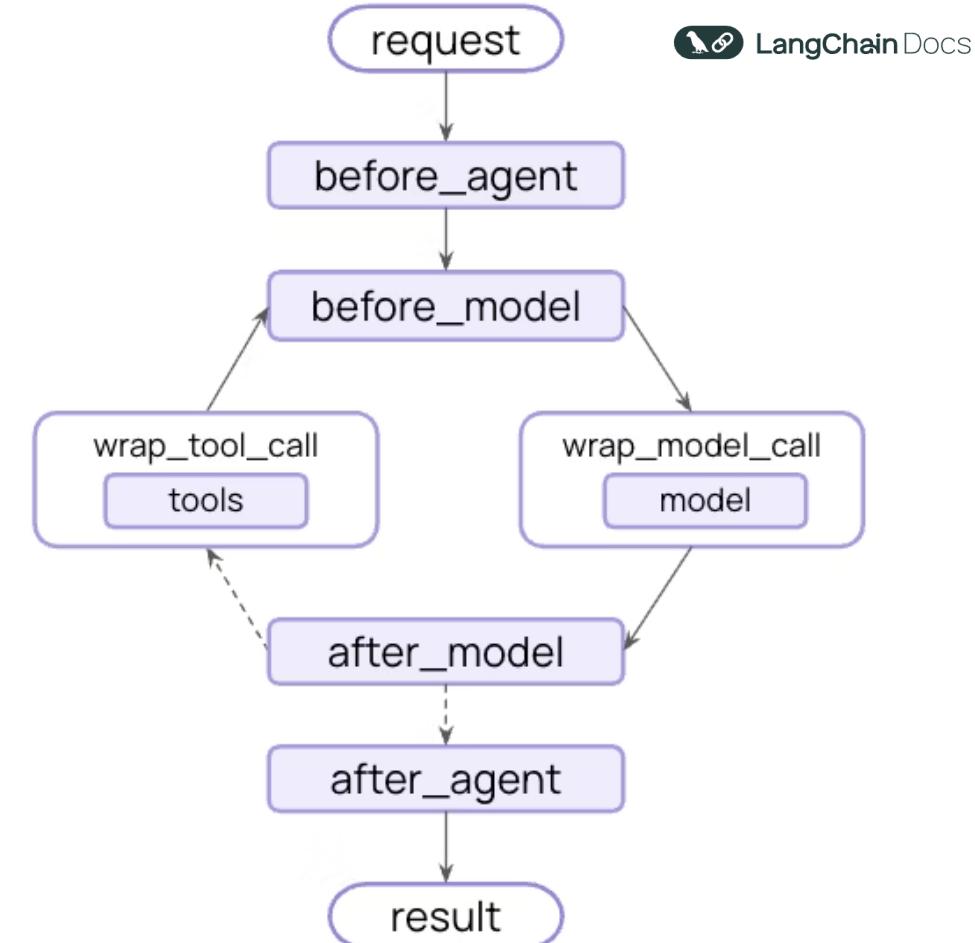


## Example: Indirect Prompt Injection



# Guardrails

- Guardrails are **explicit rules or constraints built around AI systems**, so they stay on-topic, avoid disallowed content, protect personal or sensitive data, and don't follow unintended instructions.
- **They check what the user is asking (input) and what the AI returns (output)** to detect and stop risks like hidden instructions ("prompt injections"), inappropriate or unsafe responses, or leakage of private information.
- If a guardrail detects a violation (e.g., "tell me how to build a bomb"), **the system can block the response**, reformulate it, sanitize it, or redirect the conversation; ensuring the AI behaves in a trustworthy, compliant way.



# Types of Guardrails

## Rule-Based

- **Pattern matching:** Regex for URLs, credit cards, social security numbers (e.g. Microsoft Presidio).
- **Semantic comparison:** Check similarity to known adversarial or prompt-leakage examples; block above threshold.
- **Entity removal:** Detect and mask personal data via defined entity patterns.

## Model-Based

- **AI classifiers:** e.g. Llama Guard categorizing inputs/outputs into safe/unsafe classes (A, B, C).
- **LLM-as-a-Judge:** another model reviews responses against content or compliance policy.
- **Entity removal:** Can also be learned (context-aware) using AI-based detectors.

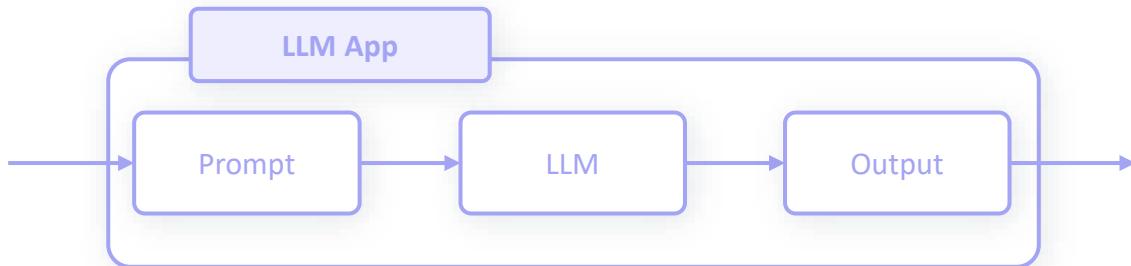


**Additional Token Consumption & Response Streaming must be considered**

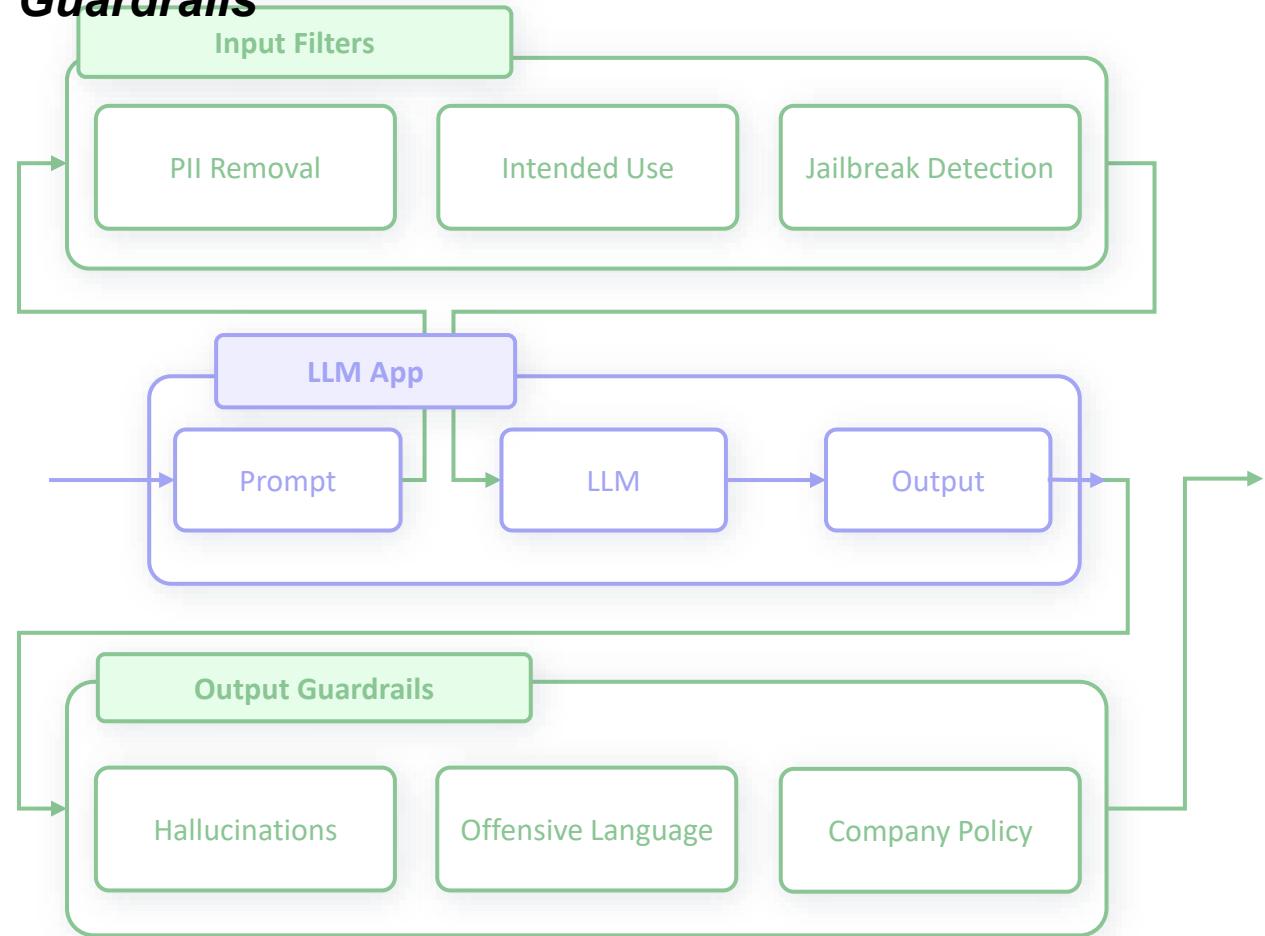


# Guardrails Example

## Without Guardrails



## With Guardrails



# What Attackers Want

Prompt Injection is not just about extracting secrets.

## System Prompt Leakage

Extract hidden instructions and security control logic

## Safety Guardrail Bypass

Generate prohibited, harmful, or policy-violating content

## Sensitive Information Disclosure

Steal PII, credentials, training data, internal documents or other context window contents

## Misinformation Induction

Force generation of false, biased, or misleading outputs

## Unauthorized Tool Execution

Invoke APIs/plugins with attacker-controlled parameters

## Resource Exhaustion

Trigger denial of service or inflate operational costs



# Realistic Defense Strategy

Since no definitive protection exists, focus on defense in depth, detection, and response

## ARCHITECTURE

Minimize attack surface through design choices

- Minimize sensitive context
- Restrict Capabilities (least privilege)
- Isolate Untrusted Data
- Enforce Structure

→ Reduce your exposure

## DETECTION

Layer multiple detection mechanisms

- Input Screening
- Output Monitoring / Structural Validation
- Behavioral Analysis
- Red Teaming & Threat Modeling

→ Each layer catches different attack patterns

## RESPONSE

Monitor, attribute, and respond to incidents

- User accountability
- Human-in-the-loop escalation
- Rate Limiting
- Automated circuit breakers

→ Make attacks attributable and costly for adversaries



# Key Takeaways

## 1 Prompt injection exploits how LLMs work

Models treat all text as potential instructions

## 2 Rule-based guardrails have limits

Pattern matching fails; semantic similarity helps partially

## 3 No complete prevention exists

Sophisticated attackers will find bypasses

## 4 Defense = Architecture + Detection + Response

Minimize exposure, layer detection, monitor, respond

