

Threat Modeling AI

What could possibly go wrong? What can we do about it?

Prof. Dr. Andrei Kucharavy

- Assistant Professor
 @ Informatics Institute of HEVS
- Co-founder
 @ HES-SO Gen Learning Center
- Cyber-Defence Campus Fellow (2020)
 “Generative ML in Cyber-Defence”
- Safety and Security
 @ Apertus Team
- “On-Premises LLMS: the Safe Way”
 @AMLD ‘24/’25
- Scientific Editor:
 LLMs in Cybersecurity (Springer)
- Contributor: OWASP & NIST PWG GenAI

Hes·so // VALAIS
WALLIS

 Hochschule für
Wirtschaft



AMLD

^ P E R T V S



Andrei Kucharavy · Octave Plancherel ·
Valentin Mulder · Alain Mermoud ·
Vincent Lenders *Editors*

Large
Language
Models in
Cybersecurity

Threats, Exposure and Mitigation

OPEN ACCESS

 Springer

Who We Are



Adrien O'Hana

- Leads the technical direction of AI audit activities
- Oversees behavioral testing and system security assessments



Gaetan Stein


- Leads business operations and client relations
- Oversees AI audit and safety activities



What is Threat Modeling?

Threat Modeling works to **identify, communicate, and understand threats and mitigations** within the context of protecting something of value.

Threat Modeling involves the **intersection of two models**: a **model of what can go wrong (threats)**, applied to a **model of the software you're building or deploying**, which is encoded in a diagram.



Why Threat Model?



- Anticipate problems when it's inexpensive to deal with them
- Communicate risk clearly to stakeholders
- Prioritize security investment where it matters most



What is YOUR Threat Model?

Four-Step Framework

1. What are you building?
2. What can go wrong with it once it's built?
3. What should you do about those things that can go wrong?
4. Did you do a decent job of analysis?

Examples

- A thief who could steal your money
- The company stakeholders who access sensitive documents
- An untrusted network
- An attacker who could steal your cookie



the grugq
@thengrugq

Following

Your threat model is not my threat model.




9:42 AM - 15 May 2017



Introduction: OWASP and GenAI Security

- OWASP (Open Worldwide Application Security Project) is a global non-profit foundation focused on improving software security.
- It regularly publishes and updates widely adopted security frameworks and best practices (e.g., OWASP Web Top 10).
- OWASP also maintains a dedicated initiative, [the OWASP GenAI Security Project](#), which focuses on identifying and mitigating security risks specific to **Large Language Models (LLMs)** and **AI agents**.



OWASP Top 10 Risk & Mitigations for LLMs and Gen AI Apps

- What if you use an open-weights model?
- What if your LLM takes no actions?
- What if your LLM has no access to sensitive information?
- What if you don't use Plugins?

👍 Useful to start conversation about LLM app security

👎 Offers little help as how to ensure it
👎 Does not necessarily apply to your system

LLM01:2025 Prompt Injection A Prompt Injection Vulnerability occurs when user prompts alter the... Read More	LLM02:2025 Sensitive Information Disclosure Sensitive information can affect both the LLM and its application... Read More	LLM03:2025 Supply Chain LLM supply chains are susceptible to various vulnerabilities, which can... Read More	LLM04:2025 Data and Model Poisoning Data poisoning occurs when pre-training, fine-tuning, or embedding data is... Read More	LLM05:2025 Improper Output Handling Improper Output Handling refers specifically to insufficient validation, sanitization, and... Read More
LLM06:2025 Excessive Agency An LLM-based system is often granted a degree of agency... Read More	LLM07:2025 System Prompt Leakage The system prompt leakage vulnerability in LLMs refers to the... Read More	LLM08:2025 Vector and Embedding Weaknesses Vectors and embeddings vulnerabilities present significant security risks in systems... Read More	LLM09:2025 Misinformation Misinformation from LLMs poses a core vulnerability for applications relying... Read More	LLM10:2025 Unbounded Consumption Unbounded Consumption refers to the process where a Large Language... Read More



NIST AI Risk Management Framework

MEASURE 1.3: Internal experts who did not serve as front-line developers for the system and/or independent assessors are involved in regular assessments and updates. Domain experts, users, AI Actors external to the team that developed or deployed the AI system, and affected communities are consulted in support of assessments as necessary per organizational risk tolerance.

Action ID	Suggested Action	GAI Risks
MS-1.3-001	Define relevant groups of interest (e.g., demographic groups, subject matter experts, experience with GAI technology) within the context of use as part of plans for gathering structured public feedback.	Human-AI Configuration; Harmful Bias and Homogenization; CBRN Information or Capabilities
MS-1.3-002	Engage in internal and external evaluations, GAI red-teaming, impact assessments, or other structured human feedback exercises in consultation with representative AI Actors with expertise and familiarity in the context of use, and/or who are representative of the populations associated with the context of use.	Human-AI Configuration; Harmful Bias and Homogenization; CBRN Information or Capabilities
MS-1.3-003	Verify those conducting structured human feedback exercises are not directly involved in system development tasks for the same GAI model.	Human-AI Configuration; Data Privacy
AI Actor Tasks: AI Deployment, AI Development, AI Impact Assessment, Affected Individuals and Communities, Domain Experts, End-Users, Operation and Monitoring, TEVV		



NIST AI Risk Management Framework

MEASURE 1.3: Internal experts who did not serve as front-line developers for the system and/or independent assessors are involved in regular assessments and updates. Domain experts, users, AI Actors external to the team that developed or deployed the AI system, and affected communities are consulted in support of assessments as necessary per organizational risk tolerance.

- 👉 Explains what needs to be done and why
- 👉 Comprehensive

- 👉 Does not explain what actually needs to be done
- 👉 Unclear how to fit into your deployment

AI Actor Tasks: AI Deployment, AI Development, AI Impact Assessment, Affected Individuals and Communities, Domain Experts, End-Users, Operation and Monitoring, TEVV



Kill Chains

- LLM systems exist within traditional cyber infrastructure.
- Attackers can **chain between** the AI kill chain and the traditional cyber kill chain; using one as an entry point to the other.





Figure 1. NVIDIA AI Kill Chain: stages of an attack on AI-powered applications

THE CYBER KILL CHAIN



👍 Useful to understand different goals of attacker at different stages

👎 Does not show interplay with other systems



MITRE ATLAS

ATLAS Matrix

The ATLAS Matrix below shows the progression of tactics used in attacks as columns from left to right, with ML techniques belonging to each tactic below. & indicates an adaption from ATT&CK. Click on the blue links to learn more about each item, or search and view ATLAS tactics and techniques using the links at the top navigation bar. View the ATLAS matrix highlighted alongside ATT&CK Enterprise techniques on the [ATLAS Navigator](#).

Filter by Maturity

Feasible Demonstrated Realized

Reconnaissance &	Resource Development &	Initial Access &	AI Model Access	Execution &	Persistence &	Privilege Escalation &	Defense Evasion &	Credential Access &	Discovery &	Lateral Movement &	Collection &	AI Attack Staging	Command and Control &	Exfiltration &	Impact &
8 techniques	12 techniques	7 techniques	4 techniques	6 techniques	8 techniques	3 techniques	11 techniques	5 techniques	9 techniques	2 techniques	4 techniques	6 techniques	2 techniques	6 techniques	8 techniques
Active Scanning & Gather RAG-Indexed Targets Gather Victim Identity Information & Search Application Repositories Search Open AI Vulnerability Analysis Search Open Technical Databases & Search Open Websites/Domains & Search Victim-Owned Websites &	Acquire Infrastructure Acquire Public AI Artifacts Develop Capabilities & Establish Accounts & LLM Prompt Crafting Obtain Capabilities & Poison Training Data Publish Hallucinated Entities Publish Poisoned Datasets Publish Poisoned Models Retrieval Content Crafting Stage Capabilities &	AI Supply Chain Compromise AI Model Inference API Access AI-Enabled Product or Service Evade AI Model Full AI Model Access Physical Environment Access Phishing & Prompt Infiltration via Public-Facing Application Valid Accounts &	AI Agent Clickbait AI Agent Tool Invocation Command and Scripting Interpreter & Deploy AI Agent Exploit Public-Facing Application & Manipulate AI Model LLM Prompt Injection Modify AI Agent Configuration User Execution &	AI Agent Context Poisoning AI Agent Tool Data Poisoning LLM Prompt Self-Replication Deploy AI Agent Manipulate AI Model Modify AI Agent Configuration Poison Training Data Prompt Infiltration via Public-Facing Application RAG Poisoning	Corrupt AI Model LLM Jailbreak Valid Accounts &	AI Agent Tool Invocation Delay Execution of LLM Instructions Evade AI Model False RAG Entry Injection Impersonation & LLM Jailbreak Unsecured Credentials &	AI Agent Tool Credential Harvesting Discover AI Agent Configuration OS Credential Dumping & RAG Credential Harvesting Discover AI Artifacts Discover AI Model Family Discover AI Model Ontology Discover AI Model Outputs Discover LLM Hallucinations Discover LLM System Information Process Discovery &	Cloud Service Discovery & Discover AI Agent Configuration Data from AI Services Data from Information Repositories & Data from Local System & Generate Deepfakes Generate Malicious Commands Manipulate AI Model Verify Attack	Phishing & Use Alternate Authentication Material &	AI Artifact Collection Craft Adversarial Data Create Proxy AI Model Exfiltration via AI Inference API Exfiltration via Cyber Means Extract LLM System Prompt LLM Data Leakage LLM Response Rendering	AI Service API Reverse Shell	Exfiltration via AI Agent Tool Invocation Data Destruction via AI Agent Tool Invocation Denial of AI Service Erode AI Model Integrity Erode Dataset Integrity Evade AI Model External Harms Spamming AI System with Chaff Data	Cost Harvesting Data Harvesting Denial of Service Erode Model Integrity Erode Dataset Integrity Evade Model External Harms Spamming AI System with Chaff Data		




MITRE ATLAS

ATLAS Matrix

The ATLAS Matrix below shows the progression of tactics used in attacks as columns from left to right, with ML techniques belonging to each tactic below. ⁸ indicates an adaption from ATT&CK. Click on the blue links to learn more about each item, or search and view ATLAS tactics and techniques using the links at the top navigation bar. View the ATLAS matrix highlighted alongside ATT&CK Enterprise techniques on the ATLAS Navigator.

Filter by Maturity


Feasible Demonstrated Realized



STRIDE over DFD

External Endpoints		
	Strengths	Weaknesses
Spoofing		VULN01: Modify System prompt (prompt injection)
Tampering		VULN02: Modify LLM parameters (temperature, length, model, etc.)
Repudiation	Proper authentication and authorization (assumed)	
Information Disclosure		VULN03: Input sensitive information to a third-party site (user behavior)
Denial of Service		
Elevation of Privilege	Proper authentication and authorization (assumed)	


LLM		
	Strengths	Weaknesses
Spoofing		VULN05: Output controlled by prompt input (unfiltered)
Tampering		VULN05: Output controlled by prompt input (unfiltered)
Repudiation		
Information Disclosure		
Denial of Service		
Elevation of Privilege		



- 👍 Specific to your application
- 👍 Explains where you can stop the attacker
- 👎 Exhausting and unsexy

MAESTRO (CSA)

"Multi-Agent Environment, Security, Threat, Risk, and Outcome"



👍 Designed for Agentic AI

👎 Overkill for simpler LLM-powered applications




STRIDE over DFD for LLM apps

Data Flow Diagrams

4 main components of a DFD

Element	Shape	Definition
Process	○	Task that receives, modifies, or redirects input to output
Data store		Permanent and temporary data storage
External entity	□	Task, entity, or data store outside of your direct control
Data-flow	↔	Data movement between processes, data stores, and external entities
Trust boundary	□	Trust zone changes as data flows through the system



STRIDE

STRIDE comes from Microsoft's security team in the 90s.

S	Spoofing Pretending to be someone/something else
T	Tampering Unauthorized modification of data
R	Repudiation Denying having performed an action
I	Information disclosure Exposing data to unauthorized parties
D	Denial of Service Making Resources unavailable
E	Elevation of Privilege Gaining unauthorized capabilities


It's a checklist of threat *types* to consider:

- Can someone fake identity?
- Modify data?
- Deny actions?
- Leak info?
- Break availability?
- Gain unauthorized access?



STRIDE for LLM apps

- Initially formulated in late 2023 (Inspired by Black Hat AI Village and Vogelsang & Majumdar's article in Andrei's LLMs in Cybersecurity book)
- Predicted several attacks & contributed to several internal threat modeling frameworks



Intersection of two models:

- How does the app work?
- What can an attacker do?

Ingress point:

- Where can the attacker enter?

Attack target:


- What is the attacker seeking to achieve?

Trust boundary:

- Where can we stop the attacker?



STRIDE for LLM apps




S	Spoofing (data authenticity) Make LLM read from the wrong database
T	Tampering (data integrity) Make LLM cite poisoned information
R	Repudiation (data origin) Make LLM cite the wrong source
I	Information disclosure (privacy) Make LLM extract private documents
D	Denial of Service (data availability) Exhaust resources via expensive queries
E	Elevation of Privilege (access rights) Make LLM run sudo shell commands




Top 5 Threats & Mitigations

Input Sanitization



MTG-1 – Input Sanitization

Do not use a prompted or fine-tuned LLM!



Common Open-Source approach:

- LLaMA-Guard (X generative)
- Prompt-Guard



Open Source:

- [Rebuff.ai](#)
- [NVIDIA NeMo Guardrails](#)
- [LLM-Guard](#)



API access:


- [Lakera Guard](#)
- [Vijil Dome](#)



Additional popular Guardrails Tools

Most LLM/Cloud APIs now offer built-in Guardrail capabilities:

Azure AI Content Safety



Amazon Bedrock
Guardrails



Checks Guardrails API is now available in alpha version in [private preview](#).



OpenAI Guardrails

Preview

Input Guardrails

Mask PII Hybrid
Detects and masks Personally Identifiable Information (PII) in text content.

Docs ↗

⌚ 500 ms

Moderation API
Blocks text flagged by moderation classifiers.

Docs ↗

⌚ 300 ms

Jailbreak LLM
Detects attempts to jailbreak LLM calls via role-playing, system prompt overrides and injections.

Docs and Evals ↗

⌚ 1000 ms

Off Topic Prompts LLM
Checks that the content stays within the defined business scope.

Docs ↗

⌚ 1000 ms

Output Guardrails

URL Filter RegEx
Blocks outputs with URLs not matching an allow list.

Contains PII Hybrid
Checks that the text does not contain personally identifiable information (PII) such as SSNs, phone numbers, credit card numbers, etc., based on configured entity types.

⌚ 10 ms

Docs ↗

⌚ 50 ms

Docs and Evals ↗

⌚ 8000 ms

Docs ↗

⌚ 1000 ms

Docs and Evals ↗

⌚ 1000 ms

Hallucination Detection OpenAI File Search
Blocks outputs with hallucinations in AI-generated text using OpenAI Responses API with file search. Validates claims against actual documents and flags potentially fabricated information.

Custom Prompt Check LLM
Block on custom moderation criteria via a text prompt.

NSFW Text LLM
Detects NSFW (Not Safe For Work) content in text, including sexual content, hate speech, violence, profanity, illegal activities, and other inappropriate material.

Custom Prompt Check LLM
Block on custom moderation criteria via a text prompt.


Docs ↗

⌚ 1000 ms



MTG-2 - Data Access Control

Do not rely on LLM alone for data access control!



Control methods of Database Security

Last Updated : 15 Dec, 2021

Database Security means keeping sensitive information safe and prevent the loss of data. Security of data base is controlled by Database Administrator (DBA).

The following are the main control measures are used to provide security of data in databases:

1. Authentication
2. Access control
3. Inference control
4. Flow control
5. Database Security applying Statistical Method
6. Encryption

These are explained as following below.

Dependencies

ML Python code
(App)

App & Dataflow
Piping

Python
Dependencies

Build dependencies

Hardware

Infrastructure

You
write
this

You
trust
this

DeepSeek AI tools impersonated by info-stealer malware on PyPI

The Hacker News

By Bill

Subscribe – Get Latest News

Home

Cyber Attacks

Vulnerabilities

Expert Insights

Contact



OpenAI Reveals Redis Bug Behind ChatGPT User Data Exposure Incident

Mar 25, 2023 · Ravie Lakshmanan


OpenAI on Friday disclosed that a bug in the Redis open source library was responsible for the exposure of other users' personal information and chat titles in the upstart's ChatGPT service earlier this week.

The [glitch](#), which came to light on March 20, 2023, enabled certain users to view brief descriptions of other users' conversations from the chat history sidebar, prompting the company to temporarily shut down the chatbot.



MTG-3 – Dependencies

Avoid version locking when you can!



Kreb's rule of online security 1:

If you didn't go looking for it, don't install it!

→ Validate dependencies

Kreb's rule of online security 2:

If you installed it, update it.

→ Weekly update rebuilds



Kreb's rule of online security 3:

If you no longer need it, remove it.

[1] <https://krebsonsecurity.com/2011/05/krebss-3-basic-rules-for-online-safety/>




External References



MTG-4 - External References

Treat external references as user inputs




Treat external references as user inputs

- Assume all external data as containing prompt injections
- Sanitize all external data as if they were user input

→ Ilama/Prompt-Guard

→ Better: referenced data control

Code Generation




Top 10 Anti-Patterns

RANK	ANTI-PATTERN	SCORE	KEY STATISTIC
#1	Dependency Risks (Slothsquatting)	24	5-21% of AI packages don't exist
#2	XSS Vulnerabilities	23	86% failure rate in AI code
#3	Hardcoded Secrets	23	Scraped within minutes of exposure
#4	SQL Injection	22	Thousands of instances in training data
#5	Authentication Failures	22	75.8% false confidence in AI auth
#6	Missing Input Validation	21	Root cause of all injection attacks
#7	Command Injection	21	CVE-2025-53773 real-world RCE
#8	Missing Rate Limiting	28	Very high frequency, easy to detect
#9	Excessive Data Exposure	28	APIs return full objects
#10	Unrestricted File Upload	28	Critical severity, enables RCE



MTG-5 - Code Generation

Treat code generated by LLMs as untrusted!



- Run LLM-generated code in sandboxed environment
 - Docker
 - chroot jail
- Block the use of libraries outside allowlist
- Block internet connections

Threat Modeling Will Take You Only So Far



*No plan survives
contact with the enemy*
- von Moltke, 1871

Real attacker are not limited by what defenders have anticipated.

Hence

We need the ability to detect them and analyze their actions.

Logging is essential



Logging: What to Do

1. Logs must be sufficiently detailed
2. Logs must not be removable
 - Deletion impossible even by admins
3. Logging must not be “off-turnable”
4. Logs must be retained for sufficiently long
 - Investigation must be able to use them
5. Logs must be treated as private information
6. Logs must be monitored
- 7. Action must be taken in case of anomalies**




Logging: How To Do

- If local software logging system
 - Attach to it
- Otherwise
 - Local Logs
 - Push to External Services
 - WandB
 - Sentry
 - Discord
 - Slack
 - Mail
 - ...



LLM Ops Lifecycle



Evaluations & Observability

Good Practices

Evaluations must evolve with the product.

Criteria are not fixed: they must follow the evolution of the model and its uses.

Benchmarks are indicators, not decisions.

They guide thinking, but they do not replace human or contextual judgment.

Your intuition (“vibe check”) becomes your first evaluation.

Perceived anomalies or weak signals during testing are already useful alerts.

Don’t overcomplicate everything with scores from 1 to 5.

Sometimes, a binary evaluation (sure / not sure, valid / not valid) is more effective.



Who we are.

Surelio.ai combines scientific, technical and operational expertise to evaluate how AI systems behave in real-world, high-stakes environments, ensuring organisations can deploy AI with confidence.

We work with private companies, industrial groups, public institutions and emerging tech organisations, always under strict confidentiality.

Meet the team



Gaetan Stein
Operations Lead



Adrien O'Hana
Technology lead



Prof. Dr. Andrei Kucharavy
Scientific Lead



Prof. Dr. Dimitri Percia David
Scientific Lead



Rémi Sabonnadière
Strategic advisor



Elliott Bertrand
Strategic advisor



Unlimtrust Campus



Thank you!

Please take 5 minutes
to provide us with
feedback:



Schedule a meeting
with Surelio.ai



HSLU Hochschule
Luzern

Hes·so VALAIS
WALLIS


ARTEFACT

EPFL
RMLD

EXOSCALE

 EXOSCALE

Swiss Cloud - European Trust

Master your stack: a sovereign cloud, built for devs by devs

Code, deploy, scale – with full sovereignty

Re-run the workshop at home

Nothing complicated: follow these 5 steps.

Scan the QR code

[Open the dedicated coupon page.](#)

Create your organization

Free Exoscale account. Use an email you check.

Request A30 GPU access

Go to Instances, click Add, select region CH-GVA-2, then choose instance type GPUA30. Click Enable to submit your access request.

100 CHF coupon already applied

Your account is credited with 100 CHF prepaid, usable on GPU.

Run your demos

If you want to go further and plan the next 90 days, contact us.



50 CHF coupon

Prefer a link? [Open the coupon page.](#)

GPU access: for governance reasons, activation is not instant. Submit your request now. We can help at any time. Contact support@exoscale.com.

Data residency EU/CH • Encryption by default • transparent egress

Managed Kubernetes and databases

Geneva • 2025